

Virtualizing Security

flyyy

Blue-Whale

1

Definitions

Some definitions before begin



2

Virtualizing development and types



3

Virtualizing realization

Qemu-kvm anylize



4

Security

How to escape from virtualizing

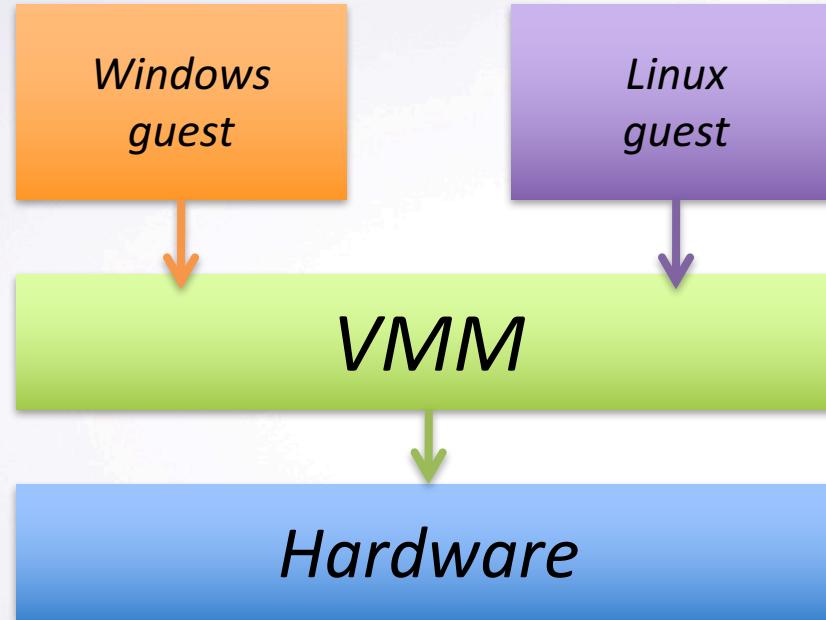


Definitions

1. VMM(Vitural Machine Monitor)and hypervisor
2. Host OS and Guest OS
3. Full-virtualization and para-virtualization

Just talk about full-virtualization.

Types

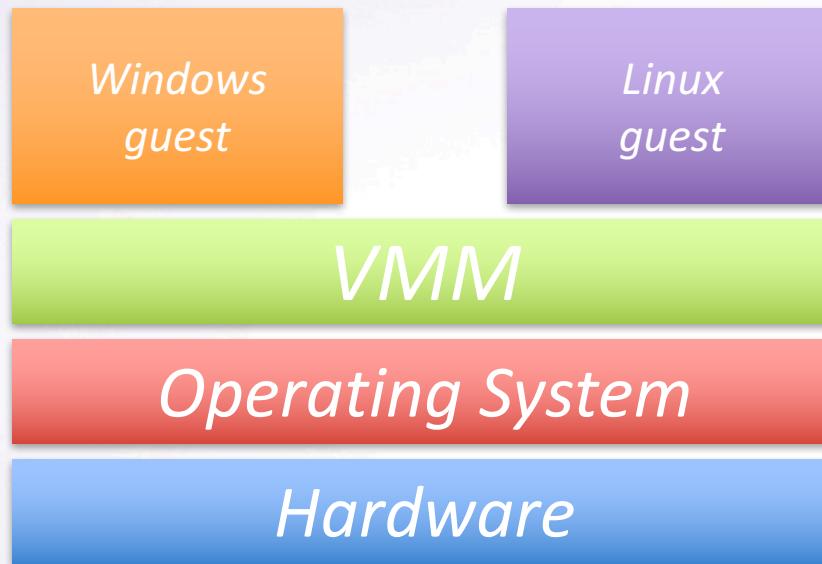


Type-I (Native)

The VMM runs directly on the hardware. It can control the hardware resources.

Examples:Vmware ESXi,
Microsoft Hyper-V.

Types



Type-II (Hosted)

The VMM runs as an application. The host OS controls the hardware resources.
Examples: VMware workstation, Qemu

Development

There used to be two ways to realize virtualization:

1. Emulation

Slow, but can realize virtualization on different architectures.

i.e. You can run mips ELF by Qemu on x86 architecture

2. Direct native execution

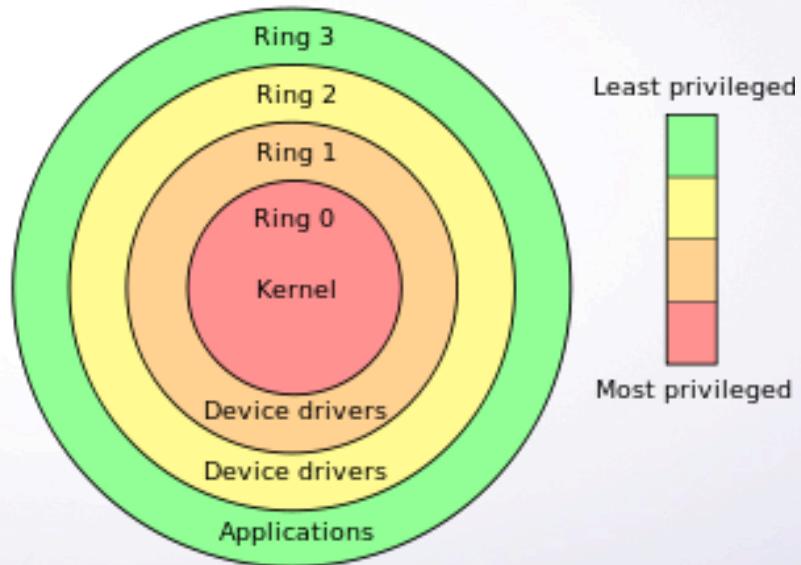
Fast, but there's a problem.

Ring0-Ring3

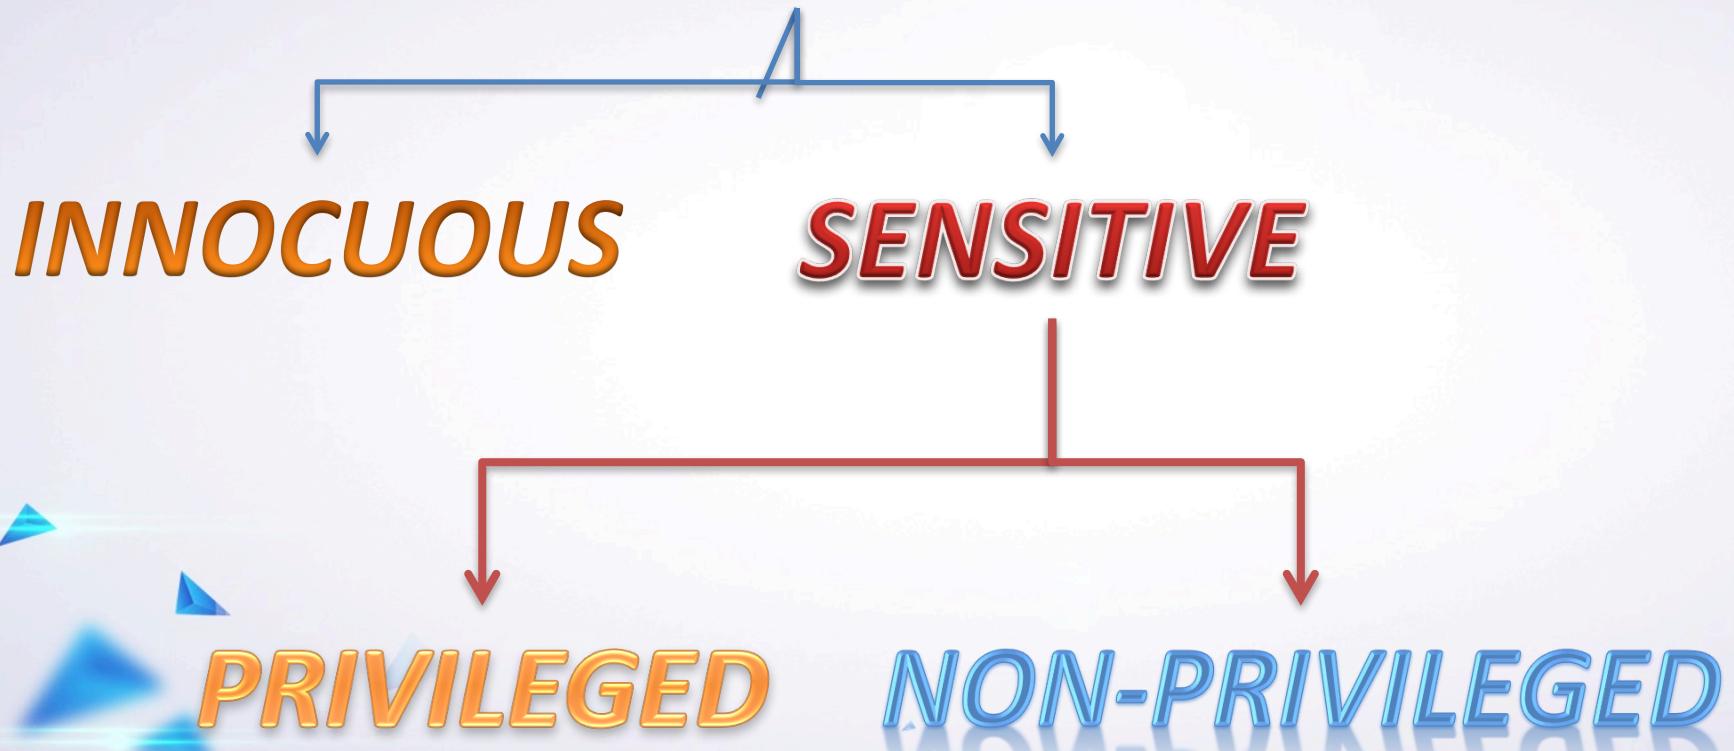
OS runs in Ring0.

User application runs in ring3.

Linux and windows only use these two levels.



INSTRUCTIONS



Sensitive instruction

It can effect system resources or behavior.

Run privileged instruction on ring3 will rise a exception.

We can catch the exception and emulate the privileged instruction.

But not all the sensitive instructions are privileged.Like popfd.

So we need to translate sensitive but not privileged instructions.

Hardware Support

Intel Virtualization Technology

New instruction : Intel Virtual Machine eXtensions (VMS)

AMD Virtualization

New instruction : AMD Secure Virtual Machine (SVM)

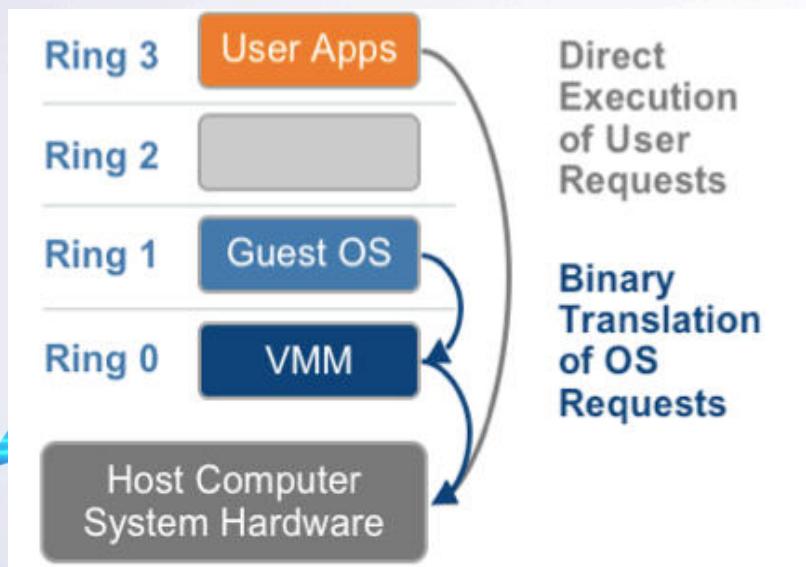
Hardware Support

VMX Root operation and VMX non-root operation.

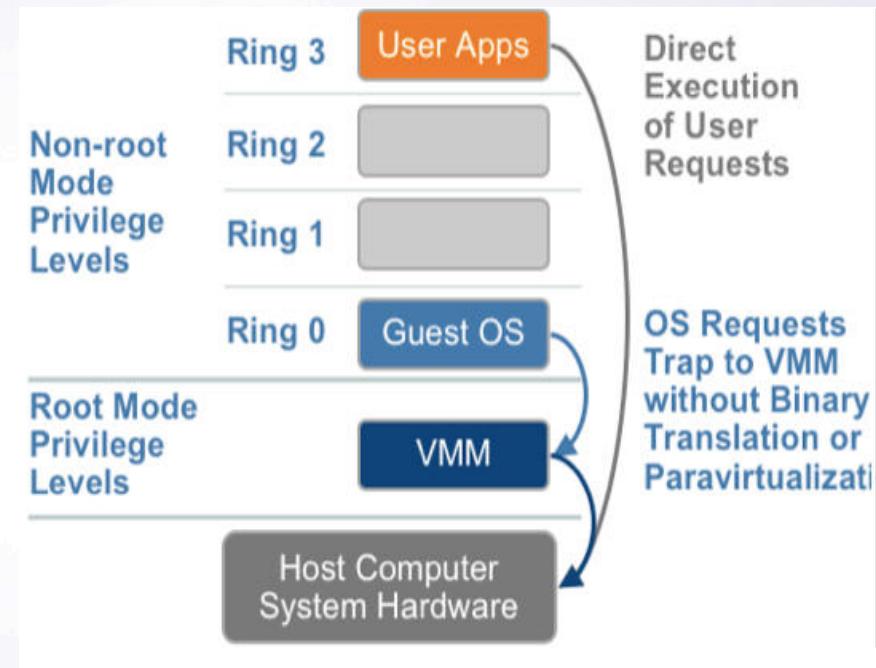
Both mode support Ring0-Ring4

Someone says VMM runs in Ring -1

Changes



Direct Execution of User Requests
Binary Translation of OS Requests



Non-root Mode Privilege Levels
Ring 3: User Apps
Ring 2: (grey)
Ring 1: (grey)
Ring 0: Guest OS
Root Mode Privilege Levels: VMM
Host Computer System Hardware

Direct Execution of User Requests
OS Requests Trap to VMM without Binary Translation or Paravirtualization

KVM

Kernel-based Virtual Machine (KVM) is a model of Linux kernel

Provide CPU virtualizing and memory virtualizing

Provide catching of guest OS' s IO

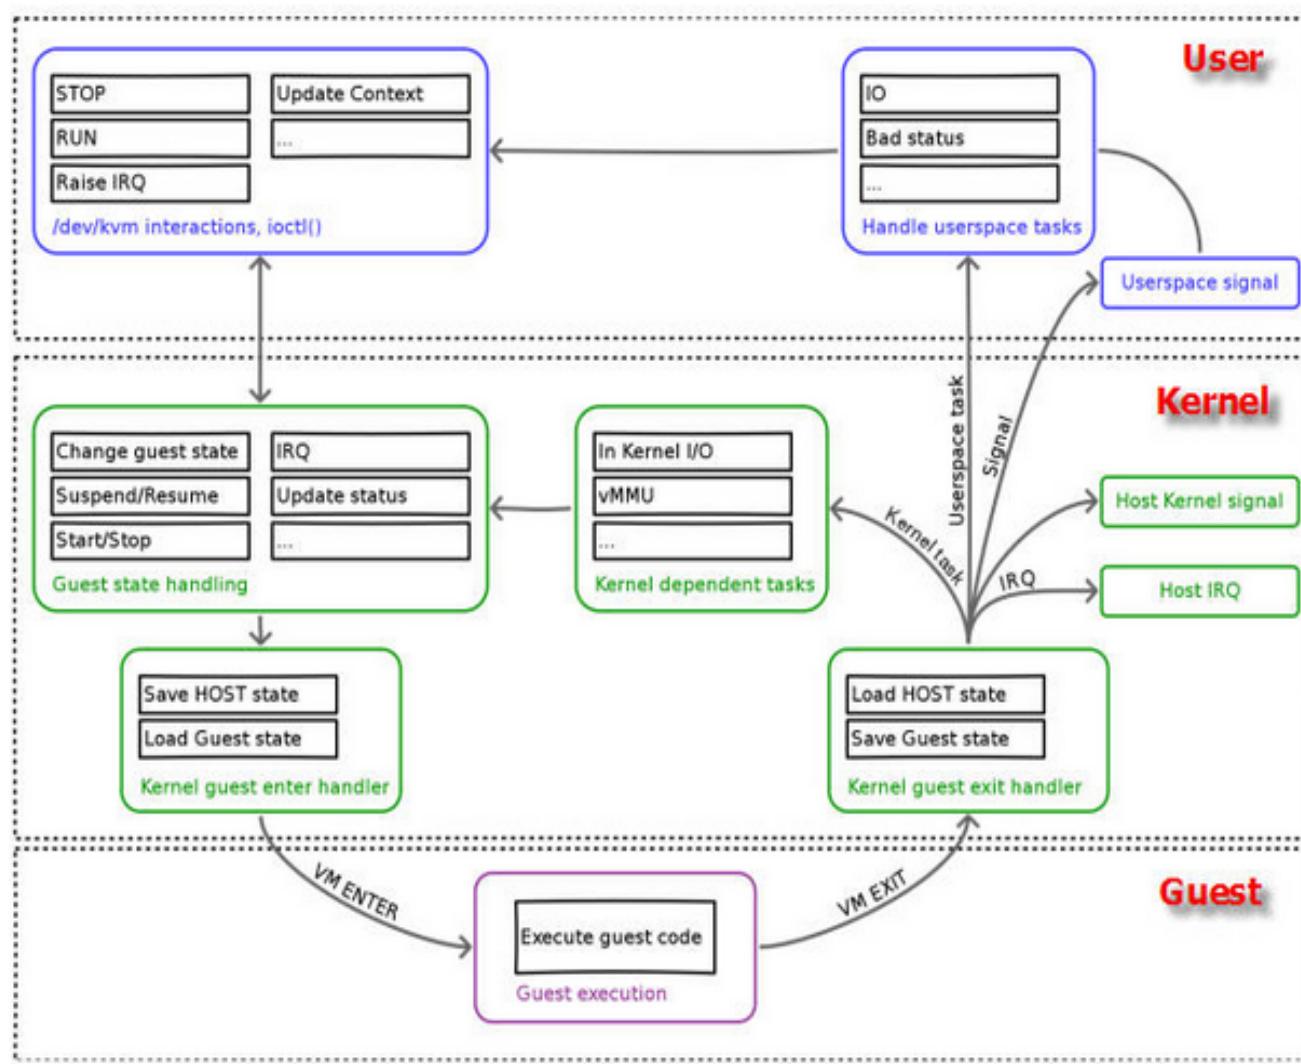
Provide API : /dev/kvm

Qemu

Qemu used to be a emulation sofaware

Now it' s used with kvm

Provided emulative hardware



```
static int (*const kvm_vmx_exit_handlers[])(struct kvm_vcpu *vcpu) = {
    [EXIT_REASON_EXCEPTION_NMI]           = handle_exception,
    [EXIT_REASON_EXTERNAL_INTERRUPT]       = handle_external_interrupt,
    [EXIT_REASON_TRIPLE_FAULT]           = handle_triple_fault,
    [EXIT_REASON_NMI_WINDOW]              = handle_nmi_window,
    [EXIT_REASON_IO_INSTRUCTION]          = handle_io,
    [EXIT_REASON_CR_ACCESS]              = handle_cr,
    [EXIT_REASON_DR_ACCESS]              = handle_dr,
    [EXIT_REASON_CPID]                   = handle_cpid,
    [EXIT_REASON_MSR_READ]               = handle_rdmsr,
    [EXIT_REASON_MSR_WRITE]              = handle_wrmsr,
    [EXIT_REASON_PENDING_INTERRUPT]       = handle_interrupt_window,
    [EXIT_REASON_HLT]                    = handle_halt,
    [EXIT_REASON_INVD]                   = handle_invd,
    [EXIT_REASON_INVLPG]                 = handle_invlpgr,
    [EXIT_REASON_RDPMC]                  = handle_rdpmc,
    [EXIT_REASON_VMCALL]                 = handle_vmcall,
    [EXIT_REASON_VMCLEAR]                = handle_vmclear,
    [EXIT_REASON_VMLAUNCH]               = handle_vmlaunch,
    [EXIT_REASON_VMPTRLD]                = handle_vmptrld,
    [EXIT_REASON_VMPTRST]                = handle_vmptrst,
    [EXIT_REASON_VMREAD]                 = handle_vmread,
    [EXIT_REASON_VMRESUME]               = handle_vmresume,
    [EXIT_REASON_VMWRITE]                = handle_vmwrite,
    [EXIT_REASON_VMOFF]                  = handle_vmoff,
    [EXIT_REASON_VMON]                   = handle_vmon,
    [EXIT_REASON_TPR_BELOW_THRESHOLD]     = handle_tpr_below_threshold,
    [EXIT_REASON_APIC_ACCESS]            = handle_apic_access,
    [EXIT_REASON_APIC_WRITE]              = handle_apic_write,
    [EXIT_REASON_EOI_INDUCED]             = handle_apic_eoi_induced,
    [EXIT_REASON_WBINVD]                 = handle_wbinvd,
    [EXIT_REASON_XSETBV]                 = handle_xsetbv,
    [EXIT_REASON_TASK_SWITCH]             = handle_task_switch,
    [EXIT_REASON_MCE_DURING_VMENTRY]      = handle_machine_check,
    [EXIT_REASON_EPT_VIOLATION]            = handle_eptViolation,
    [EXIT_REASON_EPT_MISCONFIG]            = handle_ept_misconfig,
    [EXIT_REASON_PAUSE_INSTRUCTION]       = handle_pause,
    [EXIT_REASON_MWAIT_INSTRUCTION]        = handle_mwait,
    [EXIT_REASON_MONITOR_TRAP_FLAG]       = handle_monitor_trap,
    [EXIT_REASON_MONITOR_INSTRUCTION]      = handle_monitor,
    [EXIT_REASON_INVEPT]                  = handle_invept,
    [EXIT_REASON_INVPID]                 = handle_invpid,
    [EXIT_REASON_XSAVES]                  = handle_xsaves,
    [EXIT_REASON_XRSTORS]                 = handle_xrstors,
    [EXIT_REASON_PML_FULL]                = handle_pml_full,
    [EXIT_REASON_PREEMPTION_TIMER]         = handle_preemption_timer,
};
```

Security

How to escape from a virtualizing environment?

find bugs in the code host OS executes

1. Virtualizing hardware
2. Binary Translation

Get a qemu cve

Review the source code or simple fuzz.

PCI device IO:

```
#include <sys/io.h>
iopl(3)//赋予当前程序读写IO端口的权限,也可以使用IOPERM()来临时启用
//读取
inb(port);//byte
inw(port);//word=2 bytes
inl(port);//dwords=4 bytes
//写入
outb(val,port);
outw(val,port);
outl(val,port);
```

```
#include <asm/io.h>
#include <linux/ioport.h>

long addr=ioremap(ioaddr,iomemsize);
readb(addr);
readw(addr);
readl(addr);
readq(addr);//qwords=8 btyes

writeb(val,addr);
writew(val,addr);
writel(val,addr);
writeq(val,addr);
iounmap(addr);
```

A simple bug: CVE-2017-8380

qemu scsi device : Megasas, leak memory

```
case MFI_SEQ:
    trace_megasas_mmio_writel("MFI_SEQ", val);
    /* Magic sequence to start ADP reset */
    if (adp_reset_seq[s->adp_reset] == val) {
        s->adp_reset++;
    } else {
        s->diag = 0;
    }
    if (s->adp_reset == 6) {
        s->diag = MFI_DIAG_WRITE_ENABLE;
    }
    break;
static int adp_reset_seq[] = {0x00, 0x04, 0x0b, 0x02, 0x07, 0x0d};
    ...
    ...
    ...
    case MFI_DIAG:
        retval = s->diag;
        trace_megasas_mmio_readl("MFI_DIAG", retval);
        break;
    case MFI_OSP1:
```

Try some function that can not get easily

```
qemu-system-x86_64 -m 4096 --enable-kvm -device megasas,id=bus0 -drive  
file=16GB.img,if=none,id=d0 -device scsi-disk,drive=d0,bus=bus0.0 -hda centos_i386.img
```

CTF games

Reverse : Monitor

Pwn : Qemu escape,Qemu feature(0ctf finals)

THANKS

